

# A First Introduction to MDX

This first chapter introduces the syntax and semantics of the MDX (*MultiDimensional eXpressions*) language, looking at basic queries and the language's modular nature. We will assume that you have a basic understanding of the multidimensional structures and metadata supported by the server(s) that you work with, but we won't assume that you've ever seen MDX before. This chapter introduces most of the major aspects of an MDX query and builds an important foundation for the subsequent chapters. It also introduces you to many important parts of the language. This material may bear reading slowly and perhaps more than once if you are new to the language.

## What Is MDX?

---

MDX is a language that expresses selections, calculations, and some metadata definitions against an Online Analytical Processing (OLAP) database, and provides some capabilities for specifying how query results are to be represented. Unlike some other OLAP languages, it is not a full report-formatting language, though. The results of an MDX query come back to a client program as data structures that must be processed in some way to look like a spreadsheet, a

chart, or some other form of output. This is quite similar to how SQL works with relational databases and how its related application programming interfaces (APIs) behave. As of this writing, there are several different APIs that support MDX, including Object Linking and Embedding Data Base for Online Analytical Processing (OLE DB for OLAP), ADO MD, ADOMD.Net, XMLA (*XML for Analysis*), the Hyperion Essbase C and Java APIs, and the Hyperion ADM API.

The specification for OLE DB for OLAP describes the full relationship between MDX queries and the data structures that convey the queried information back to the client program. This chapter focuses more on the more logic-related side—what queries are asking for—rather than the programming-oriented aspect of how the results are returned.

## Query Basics

We will start by looking at simple MDX queries and build on them in small steps. Throughout this chapter, we will mix descriptions of the abstract properties of a query in with concrete examples to build up a more comprehensive picture of MDX.

Imagine a very simple cube, with dimensions of time, sales geography, and measures. The cube is called Sales. Let's say we want to look at a grid of numbers that has the Massachusetts sales and costs for the first two quarters of 2005. MDX queries result in grids of cells holding data values. The grid can have two dimensions, like a spreadsheet or table, but it can also have one, three, or more. (It can also have zero; we'll talk about that in the section "Queries.") The grid we want to see is shown in Figure 1-1.

The MDX query in the following example would specify the cells we want to see in Figure 1-1:

```
SELECT
{ [Measures].[Dollar Sales], [Measures].[Unit Sales] }
on columns,
{ [Time].[Q1, 2005], [Time].[Q2, 2005] }
on rows
FROM [Sales]
WHERE ([Customer].[MA])
```


	Dollar Sales	Unit Sales
Q1, 2005	96,949.10	3,866
Q2, 2005	104,510.20	4,125

**Figure 1-1** A very simple data grid to retrieve.

In the query, `SELECT`, `FROM`, and `WHERE` are keywords that denote different parts of the query. The result of an MDX query is itself a grid, essentially another cube. You lay out the dimensions of the cube you are querying onto the *axes* of the result; this query refers to two of them by the names *rows* and *columns*. In MDX terminology, an *axis* is an edge or dimension of the query result. Referring to “axis” rather than “dimension” makes it simpler to distinguish the dimensions of the cube being queried from the dimensions of the results. Furthermore, each axis can hold a combination of multiple cube dimensions.

You may guess some generalizations immediately from this example. Let’s break this simple query down into pieces:

1. The `SELECT` keyword starts the clause that specifies what you want to retrieve.
2. The `ON` keyword is used with an axis name to specify where dimensions from your database are displayed. This example query puts measures on the columns axis and time periods on the row axis.
3. MDX uses curly braces, { and }, to enclose a set of elements from a particular dimension or set of dimensions. Our simple query has only one dimension on each of the two axes of the query (the measures dimension and the time dimension). We can separate different elements with commas (,). Element names may be enclosed by [ ] characters, and may have multiple parts separated by dot (.) characters.
4. In an MDX query, you specify how dimensions from your database lay out onto axes of your result. This query lays out measures on columns and time periods on rows. Each query may have a different number of result axes. The first three axes have the names “columns,” “rows,” and “pages” to conceptually match a typical printed report. (You can refer to them in other ways, as you will see in the “Axis Numbering and Ordering” section.) Although this simple query does not show more than one dimension on a result axis, when more than one dimension maps to a result axis, each cell slot on the axis is related to a combination of one member from each of the mapped dimensions.
5. The `FROM` clause in an MDX query names the cube from which the data is being queried. This is similar to the `FROM` clause of Structured Query Language (SQL) that specifies the tables from which data is being queried.
6. The `WHERE` clause provides a place to specify members for other cube dimensions that don’t appear in the columns or rows (or other axes). If you don’t specify a member for some dimension, then MDX will make use of some default. (We will ignore the parentheses in this query until later on when we discuss tuples.) The use of `WHERE` is optional.



Once the database has determined the cells of the query result, it fills them with data from the cube being queried.

MDX shares the keywords `SELECT`, `FROM`, and `WHERE` with SQL. Based on substantial experience teaching MDX, please take this advice: If you are familiar with SQL and its use of `SELECT`, `FROM`, and `WHERE`, do your best to forget about SQL when learning MDX. The meanings and semantics are quite different, and trying to apply SQL concepts to MDX will most likely create confusion.

Let's look at another example. As far as MDX is concerned, every dimension is like every other dimension, even the measures dimension. To generate the result grid shown in Figure 1-2, with sales shown over all three quarters across the columns and for two states down the rows, we can use the following query:

```
SELECT
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
on columns,
{ [Customer].[MA], [Customer].[CT] }
on rows
FROM Sales
WHERE ( [Measures].[Dollar Sales] )
```

As you can see, the orientation of time, location, and measures has been switched just by listing time on columns, customers on rows, and measures in the `WHERE` section.

MDX has a number of other relevant aspects that the remainder of this chapter will describe:

- The MDX data model
- Simple MDX query construction
- Parts of MDX queries

You can skip ahead to the section “Simple MDX Construction,” later in this chapter, if you want to jump right into learning more MDX vocabulary. In the discussion of the first two topics—the MDX data model and how metadata entities are named—we will present a lot of detail that may seem mundane or boring. However, the details are important for almost every query, and gaining a deep understanding of MDX demands mastery of the basics. If you do jump ahead and later find yourself getting lost in the syntax, come back to this chapter and get a refresher on the basics before moving on to the next level.

	Q1, 2005	Q2, 2005	Q3, 2005
MA	96,949.10	104,510.20	91,025.00
CT	12,688.40	24,660.70	16,643.90

**Figure 1-2** A second simple result grid.

## Axis Framework: Names and Numbering

You signify that you are putting members onto columns or rows or other axes of a query result by using the “on columns/rows/pages/etc” syntax. The axis designations can come in any order; the last query would have been just as valid and would mean exactly the same if it had been phrased as follows:

```
SELECT
{[Customer].[MA], [Customer].[CT] }
on rows,
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
on columns
FROM Sales
WHERE ( [Measures].[Dollar Sales] )
```

You can also use axis numbers to specify the axis in the query, by stating, for example:

```
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
on axis(0),
{[Customer].[MA], [Customer].[CT] }
on axis(1)
```


The phrase `on axis (n)` indicates that those members should be put on the axis numbered *n*. The names used so far are synonyms for these numbers:

0	Columns
1	Rows
2	Pages
3	Chapters
4	Sections

For axes beyond `axis (4)`, you must use the numbers, because there are no names for them. You can also mix and match names and numbers in a query:

```
SELECT
{[Customer].[MA], [Customer].[CT] }
on rows,
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
on axis(0)
FROM Sales
WHERE ( [Measures].[Dollar Sales] )
```

However, a query that uses axis 1 must also use axis 0, and a query that uses axis 2 must also use axes 1 and 0. You cannot skip an axis in a query—you’ll get an error. The following won’t work because it skips axis 1:



```
SELECT
{[Customer].[MA], [Customer].[CT] }
on axis(2),
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
on axis(0)
FROM Sales
WHERE ( [Measures].[Dollar Sales] )
```

If you think of the result of an MDX query as a hypercube (or subcube), and each axis as a dimension of that hypercube, then you can see that skipping an axis would be problematic.

## Case Sensitivity and Layout

---

MDX is neither case-sensitive nor line-oriented. For example, both of the following fragments are equally valid:


```
SELECT
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
ON COLUMNS, ...
```

```
select
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q3, 2005] }
on columns, ...
```

You can mix capitalization of letters as you want. For language keywords, capitalization doesn't matter. For names, capitalization may or may not matter to the MDX provider—it usually doesn't for Analysis Services or Essbase, although there are some cases that we will cover later where Essbase does care about capitalization. It also doesn't matter where you put spaces and line ends. You can rewrite the prior fragments as the following:

```
SELECT {
    [Time].[Q1, 2005]
, [Time].[Q2, 2005]
, [Time].[Q3, 2005]
}
ON
COLUMNS,
```

This gives you some flexibility when you are generating MDX by code as well as when you want to arrange it for readability. We strongly recommend using plenty of indenting when writing MDX queries, for both you and anyone else who has to understand your work.



**TIP** In this book, we will frequently use white space and an outline-like style in MDX examples. Our experience is that this helps make code more readable and much easier to edit, debug, and evolve. This follows the same principles as development in any other programming language. In order to conserve page space, however, we will sometimes compress the examples a bit more than we would prefer.

Additionally, MDX makes heavy use of parentheses, brackets, and braces. These can look fairly similar in common fonts like Arial. In whatever environment you write MDX, if you can use a monospace font such as Lucida Console or Courier, you will have an easier time reviewing the MDX you have written.

## Simple MDX Construction

Let's build up a little vocabulary and look at some of the very simplest and most frequently used operators and functions. We will introduce them here and describe how they are often used. (More complete and detailed descriptions can be found in Appendix A, which is a reference on MDX functions and operators.) The ones we will look at here are

- Comma (,) and colon (:)
- .Members
- .Children and Descendants()

We will use the term *set* when discussing these functions and operators. A set is a rather important concept in MDX, and will be described in detail in the section "The MDX Data Model: Tuples and Sets." Until then, we will be a little informal with our use.

### Comma (,) and Colon (:)

We have already seen the comma operator used to construct sets; let's talk about this more here. We can assemble a set by enumerating its components and separating them with commas, as in:

```
{ [Time].[January 2005], [Time].[February 2005],  
  [Time].[March 2005] }
```

This expression results in a set that holds the first three months of 2005.

**EXPRESSIONS: THE “X” IN MDX**

**MDX has a lot of little primitive “moving parts” that can be put together in many ways. Mastering MDX requires building up both your vocabulary of parts and techniques for assembling them. In this discussion (and in many to follow), we will show the operators in use both within whole queries and as fragments of MDX. We encourage you to begin trying to understand what MDX provides in terms of building blocks as well as finished products.**

At every level in every dimension, the members of that level are arranged in a particular order (usually by member key or by name). When it makes sense, you can specify a set as a range of members in that order by listing two members from the same level as endpoints and putting a colon between them to mean “These members and every member between them.” (This is similar to the syntax used to specify ranges of cells in spreadsheets like Excel.) For example, the following query, whose results are shown in Figure 1-3, selects the months from September 2004 through March 2005 and the product categories from Tools through Home Audio for customers in Lubbock, and Unit Sales:

```
SELECT
{ [Time].[Sep,2004] : [Time].[Mar,2005] } on columns,
{ [Product].[Tools] : [Product].[Home Audio] } on rows
FROM [Sales]
WHERE ([Customer].[Lubbock, TX], [Measures].[Unit Sales])
```

Sets are usually built like this when the database-defined ordering corresponds to a useful real-world ordering, such as with time. The colon takes two members on the same level as its endpoints. It is not an error to put the same member on both sides of the colon; you will just get a range of one member.

The comma operator can be used anywhere within a set specification to add subsets to an overall set. For example, the following creates a set of the first three and last three months of 2001:

```
{ { [Time].[January-2001] : [Time].[March-2001] } ,
{ [Time].[October-2001] : [Time].[December-2001] } }
```

whereas the following creates a set of 2001 and its first three months:

```
{ [Time].[2001], { [Time].[January-2001] : [Time].[March-2001] } }
```



	Sep, 2004	Oct, 2004	Nov, 2004	Dec, 2004	Jan, 2005	Feb, 2005	Mar, 2005
Tools		79	120	236			62
Toys	28	44	159	292	48	51	35
Auto Electronics		54		47			
Computers, Peripherals	48	72	204	325	67	74	53
Digital Photography							
Home Audio							

**Figure 1-3** Result of query using the colon (:) operator.

When subsets are concatenated by commas, the order in which the commas join them is the order in which they are returned.

## PRODUCT DIFFERENCES

### ANALYSIS SERVICES

**In Analysis Services, regardless of whether the member on the left of the colon is before or after the member on the right in terms of database ordering, the set that results will have all members between the two in database ordering. That is, the following two return exactly the same set:**

```
{[Time].[April-2001] : [Time].[January-2001]}
{[Time].[January-2001] : [Time].[April-2001]}
```

### ESSBASE

**In Essbase, the set starts with the left member and runs through the right member. The following returns the ordering April, March, February, January:**

```
{[Time].[April-2001] : [Time].[January-2001]}
```

**The members must also be at the same level; you can use the `MemberRange()` function to get a range from a common generation.**

## .Members

Getting the set of members for a dimension, hierarchy, or level is very common both for retrievals and as a starting point for further operations. The `.Members` operator takes a dimension, hierarchy, or level on its left-hand side, and returns a set of all members associated with that metadata scope. For example, `[Customer].Members` results in the set of all customers, whereas `[Product].[Product Category].Members` returns all members of the Product Category level in the Product dimension. For example, the query

```
SELECT
{ [Scenario].Members } on columns,
{ [Store].Members } on rows
FROM Budgeting
```

lays out all members of the scenario dimension across the columns and all members of the store dimension down the rows.



When a client uses `.Members` (or other metadata functions that return the set of members associated with some metadata element), neither Analysis Services nor Essbase will include any calculated members in the returned set. (Calculated members are a means for introducing calculations onto other data in the database, and are discussed starting in Chapter 2.) This means that the preceding request for `[Scenario].Members`, as written, will not return any calculated members in the scenario dimension. We can always ask for them by name, however, and Analysis Services provides the `AddCalculatedMembers()` and `.AllMembers` functions to add them into a set. See Appendix A for details on these functions.

## Getting the Children of a Member with `.Children`

Another kind of selection that is very frequent is to get the children of a member. We may want to do this to implement a drill-down operation, or simply to conveniently obtain a range of members based on a common parent. MDX provides a `.Children` function that will do this for us. The following query selects both the `[Product].[Tools]` member and its children on the rows of a report, which is shown in Figure 1-4:

```
SELECT
{ [Time].[Q3, 2005].Children }
on columns,
{ [Product].[Tools], [Product].[Tools].Children }
on rows
FROM Sales
WHERE ([Customer].[TX], [Measures].[Unit Sales])
```

You can get the children of any member that has children using this function. If you request the children of a leaf-level member, you'll get an empty set of members back. We'll talk more about what this means later on.



**PRODUCT-SPECIFIC ASPECTS OF .MEMBERS****ANALYSIS SERVICES 2000 AND 2005**

In Analysis Services, hierarchies behave like dimensions, and multiple dimensions have different sets of members. As a result, in an Analysis Services dimension that contains multiple hierarchies, you cannot simply request `[Dimension].Members`. If you do, it will complain of an unknown dimension. For example, given a logical `[Time]` dimension that contains two hierarchies, `[Time].[Fiscal]` and `[Time].[Calendar]`, a client taking metadata from OLE DB for OLAP will see one time dimension. However, the expression `[Time].Members` will result in an error rather than result in all members on all time hierarchies. To obtain a set of members, the client must request either `[Time].[Fiscal].Members` or `[Time].[Calendar].Members`. If the dimension has only one hierarchy and that hierarchy does not have an explicit name, then `[Dimension].Members` will work. For example, if time has only one hierarchy, then `[Time].Members` will work.

**ESSBASE**

In Essbase 9, when you request all members of some scope with `.Members`, all copies of shared members are returned in the result set. There is no way to strip shared members from a set. Therefore, you may need to build awareness of where shared members exist into your MDX, for example using `Descendants()` (described later) instead of `.Members`.

	July, 2005	Aug, 2005	Sep, 2005
Tools	484	554	319
Bench Power Tools	42	133	94
Compressors, Air Tools			51
Electrical Shop	107	118	33
Garage Door Openers	57	46	53
Hand Tools, Carpentry		55	
Hand Tools, General Purpo	138	164	31
Mechanics Tools	88	38	57
Portable Power Tools	52		
Power Tool Accessories			

**Figure 1-4** Result of query using `.Children`.

## Getting the Descendants of a Member with `Descendants()`

To request members further away than immediate children, or for a more generalized search below in the hierarchy, you use `Descendants()`. We will skim the surface of this function because it has too many options to go into here; you can refer to a complete reference for it in Appendix A.

Syntax:

```
Descendants (member [, [ level ] [, flag]] )
```

`Descendants()` returns the members below *member* related to the *level* or generation, with a number of selection options based on the optional *flag*. The choices for *flag* are


```
SELF  
BEFORE  
AFTER  
SELF_AND_BEFORE  
SELF_AND_AFTER  
SELF_BEFORE_AFTER  
LEAVES
```

`SELF` refers just to the members at *level*, and this is the most frequent option. For example, the following selects the months of 2005, which is shown in Figure 1-5:

```
SELECT  
{ [Product].[Tools], [Product].[Toys] } ON COLUMNS,  
Descendants (  
    [Time].[2005],  
    [Time].[Month],  
    SELF  
)  
ON ROWS  
FROM Sales  
WHERE [Measures].[Dollar Sales]
```

Since `SELF` is so frequently used, it is the default option if you omit the flag from your code. For example, `Descendants ([Time].[2005], [Time].[Month])` will also return the list of months in 2005 even though `SELF` is not explicitly included in the query.

The other flags that include `SELF` refer to levels around the level you reference. `SELF_AND_BEFORE` means to return all members between the `SELF` and “before” that level, which means all of them up to and including the starting member. For example, the following picks out [2005] and all quarters and months within it, and the results are shown in Figure 1-6:



```

SELECT
{ [Product].[Tools], [Product].[Toys] } ON COLUMNS,
Descendants (
    [Time].[2005],
    [Time].[Month],
    SELF_AND_BEFORE
)
ON ROWS
FROM Sales
WHERE [Measures].[Dollar Sales]

```

	Tools	Toys
Jan, 2005	59,722.40	49,948.20
Feb, 2005	65,604.10	42,885.40
Mar, 2005	57,715.50	56,601.70
Apr, 2005	64,179.90	51,794.40
May, 2005	68,152.60	62,135.70
Jun, 2005	67,476.70	55,582.90
Jul, 2005	71,997.90	50,111.80
Aug, 2005	71,411.90	48,965.30
Sep, 2005	58,979.60	52,532.90
Oct, 2005	77,720.10	58,969.60
Nov, 2005	196,946.70	147,854.50
Dec, 2005	223,948.60	171,600.20

**Figure 1-5** Result of using Descendants( , , SELF).

	Tools	Toys
2005	1,083,855.90	848,982.70
Q1, 2005	183,042.90	149,435.30
Jan, 2005	59,722.40	49,948.20
Feb, 2005	65,604.10	42,885.40
Mar, 2005	57,715.50	56,601.70
Q2, 2005	199,809.20	169,513.10
Apr, 2005	64,179.90	51,794.40
May, 2005	68,152.60	62,135.70
Jun, 2005	67,476.70	55,582.90
Q3, 2005	202,389.30	151,610.00
Jul, 2005	71,997.90	50,111.80
Aug, 2005	71,411.90	48,965.30
Sep, 2005	58,979.60	52,632.90

**Figure 1-6** Result of using Descendants ( , , SELF\_AND\_BEFORE).

`SELF_AND_AFTER` means to return all members from the level referenced down to the leaf level. `SELF_BEFORE_AFTER` means to return the entire subtree of the hierarchy down to the leaf level, starting with the given *member*. Note that this is the same as just saying `Descendants ([Time].[2005])`, since leaving off both the flag and the level means to take all descendants.

Note also that you can refer to descendants relative to a specific level, or you can refer to descendants relative to some depth away from the given member by using a number instead of a reference to a level. There are several more useful options; once again, you might be well served by looking at the reference in Appendix A.

**NOTE** In Essbase, you can refer to either generations or levels in the *level* argument.

---


## Removing Empty Slices from Query Results

---

In a multidimensional space, data is frequently sparse. For example, not all products are sold in all stores at all times to all customers, so a query for some particular set of products, time periods, and customers may return some slices along one or more axes that are entirely empty. You can be pretty sure that any large retailer won't be selling snow shovels in Hawaii in December but will have a hard time keeping them in stock in Chicago. A user may want to see the empty slices, because the empty slice may tell him or her about a complete absence of activity, but the user may also just want the empty slices treated as irrelevant and have them removed from the report. Let's take as an example the following query, the results of which are shown in Figure 1-7:

```
SELECT
{ [Time].[Jan,2005], [Time].[Feb,2005] }
ON COLUMNS ,
{ [Product].[Toys],
  [Product].[Toys].Children
}
ON ROWS
FROM Sales
WHERE ([Measures].[Dollar Sales],[Customer].[TX])
```

For some reason, we haven't sold any Dolls, Electronics, Games, or Musical toys during these two months anywhere in Texas. If we only want to see the products that we did sell within the report's time frame, MDX provides a way to remove the entirely empty slices from the query result, by using the `NON EMPTY` keywords. You can see it in use in the following query, whose results are shown in Figure 1-8:



```

SELECT
{ [Time].[Jan, 2005], [Time].[Feb, 2005] }
ON COLUMNS,
NON EMPTY
{ [Product].[Toys],
  [Product].[Toys].Children
}
ON ROWS
FROM Sales
WHERE ([Measures].[Dollar Sales], [Customer].[TX])

```

All we need to do is to add the **NON EMPTY** keywords at the beginning of our request for the axis. In this case, it is after the comma that comes after **ON COLUMNS**, and before the **{** that starts our request for products. The query is evaluated and after all the data for the columns is known, the rows where at least one column has data in it are returned.

	Jan, 2005	Feb, 2005
Toys	6,950.00	7,666.20
Action Figures	747.20	421.50
Arts & Crafts	2,499.80	2,135.90
Cars & Trucks		1,078.40
Construction	3,002.90	982.00
Dolls		
Educational	700.10	597.10
Electronics		
Games		
Musical		
Radio Controlled		2,451.30

**Figure 1-7** Query with empty slices on rows.

	Jan, 2005	Feb, 2005
Toys	6,950.00	7,666.20
Action Figures	747.20	421.50
Arts & Crafts	2,499.80	2,135.90
Cars & Trucks		1,078.40
Construction	3,002.90	982.00
Educational	700.10	597.10
Radio Controlled		2,451.30

**Figure 1-8** Empty slices removed with **NON EMPTY**.

NON EMPTY works on any axis, and with any dimensions and tuples. The query can be flipped around so that you can see this:

```
SELECT
{ [Time].[Jan, 2005], [Time].[Feb, 2005] }
ON ROWS ,
NON EMPTY
{ [Product].[Toys],
  [Product].[Toys].Children
} ON COLUMNS
FROM Sales
WHERE [Customer].[TX]
```

There are other functions that you can make use of to remove members with empty data from your results, but NON EMPTY is perfectly suited for this.

**NOTE** Analysis Services 2005 introduces a HAVING clause to queries, which is similar in effect to NON EMPTY but is more flexible. It is discussed in Chapter 4.

---

## Comments in MDX


To facilitate documentation, three variations of comment syntax are allowed in MDX, which should suit a variety of styles. (As with other programming, comments can also be used to exclude sections of code for debugging, too.)

The first style uses the symbols /\* and \*/ to delimit a comment. All characters between the /\* and the \*/ are ignored by the MDX parsing machinery. This enables comments to be placed within a line and for the comment to span lines. The following is an example of this style of comment:

```
SELECT /* Put products
on columns */ [Product].Members
on columns FROM Cube
```

The second style uses the // symbol to begin a comment, and the comment will extend to the end of the line. The following is an example:

```
SELECT // Put products on columns
[Product].Members
on columns FROM Cube
```





The third style is like the second style but uses a pair of dashes (—) to begin a comment. The following is an example:

```
SELECT — Put products on columns
[Product].Members
on columns FROM Cube
```

Comments can be placed anywhere white space could be used. For example, `[Product]./* whole dimension */ Members` will work just fine. Comments can be used in queries or expressions. Make sure that you don't use them directly inside a member name, though!

If you use comments to selectively include and exclude portions of a query or expression (for example, while debugging it), keep in mind that the `/* */` comments cannot be nested. That is,

```
/* /* comment */ */
```

is not a valid comment, whereas

```
/* /* comment */
```

is a valid comment. In the first of these two examples, the first `*/` ends the overall comment, so the second `*/` is a token that the MDX parser will try to parse (and fail on).

## The MDX Data Model: Tuples and Sets

MDX uses a data model that starts with cubes, dimensions, and members, but it is richer and somewhat more complex. Understanding this data model is the key to moving beyond the basics and unlocking powerful analyses. In this section, we will explore the key elements of this model: *tuples* and *sets*. Tuples and sets are very similar to members and dimensions, but are more generalized and more flexible.

**TIP** Understanding tuples and how to use them in the various components of MDX is sometimes the most difficult part of learning MDX. You may want to revisit this section again before proceeding with the more advanced content later in this book.

## Tuples

A *tuple* is a combination of members from one or more dimensions. It is essentially a multidimensional member. A single member is a simple tuple (for example, [Time].[Jun, 2005]). When a tuple has more than one dimension, it has only one member from each dimension. To compose a tuple with more than one dimension, you must wrap the members in parentheses, for example:

```
([Customer].[Chicago, IL], [Time].[Jan, 2005])
```

A tuple can stand for a slice of a cube, where the cube is sliced by each member in the tuple. A tuple can also be a data object that you can manipulate on its own in MDX.

We can involve tuples that we build using this syntax directly in queries. For example, consider the following, the results of which are shown in Figure 1-9:

```
SELECT
{ ( [Time].[2005], [Measures].[Dollar Sales] ),
  ( [Time].[Feb, 2005], [Measures].[Unit Sales] )
}
ON COLUMNS ,
{ [Product].[Tools], [Product].[Toys] } ON ROWS
FROM [Sales]
```

Our result was asymmetric as far as the combinations of time and measures that get returned. This gives us a great deal of fine control over the combinations of members and cells that will come back from a query.


You can always put a single member within parentheses, but it's not required if the tuple is defined by just that member. However, the following is not valid because it has two time members in it:

```
([Customer].[Chicago, IL], [Time].[Jan, 2005], [Time].[Feb, 2005])
```

Also, you can't compose an empty tuple. () is not a valid tuple. (Analysis Services 2005 has the concept of a null member that you can specify using null, so you can create a tuple by writing (null) or (null, null), but this appears to be mostly useful when using stored procedures as discussed in Chapter 10.)

In addition to composing tuples by explicitly creating them in code, a number of functions will return tuples.

The "dimensionality" of a tuple refers to the set of dimensions whose members compose it. The order in which dimensions appear in a tuple is an important part of a tuple's dimensionality. Any and all dimensions can be part of a tuple, including members of the measures dimension.



	2005	Feb, 2005
	Dollar Sales	Unit Sales
Tools	1,083,855.90	2,621
Toys	848,982.70	1,695

**Figure 1-9** Simple example of tuples on columns.

Although you can compose a tuple from members wrapped by `()`, you cannot use syntax to compose a tuple from tuples. That is, you can build up a tuple by saying:

```
(
  [Time].[2004],
  [Customer].[Chicago, IL],
  [Product].[Tools]
)
```

but not:

```
(
  [Time].[2004],
  (
    [Customer].[Chicago, IL],
    [Product].[Tools]
  )
)
```

(Using functions that let you compose tuples in other ways, you can build tuples from other tuples. We'll talk about that later. But not by putting their names together like this.)

In calculations and queries, MDX identifies cells based on tuples. Conceptually, each cell value is identified by a tuple composed of one member from each dimension in the cube. (This is sort of like a spreadsheet, in which Sheet 1, Column B, Row 22 identifies a cell.) In a query, some of these members' dimensions may be placed on rows, others on columns, others on pages, and yet others on the query slicer. However, the intersection of two or more tuples is yet another tuple, so combining them all together yields a cell in the end. The tuple `([Product].[Leather Jackets], [Time].[June-2005], [Store].[Fifth Avenue NYC], [Measures].[Dollar Sales])` may completely define a cell with a value of \$13,000.

Although a tuple either refers to a combination of members, when it is used in an expression where a number or string might be used, the default behavior is to reference the value in the cell that the tuple specifies. This is a little point that makes a big difference when trying to understand what some of the MDX functions do. How these references play out is the subject of Chapter 4.

In addition to creating tuples by explicitly coding them, a few functions also return tuples. We will describe them later in this book.

## Sets

A set is simply an ordered collection of tuples. A set may have more than one tuple, have only one tuple, or be empty. Unlike a mathematical set, an MDX set may contain the same tuple more than once, and ordering is significant in an MDX set. Although sets might be better called “collections” or “sequences,” we are stuck with “set” for now. Depending on the context in which a set is used, it either refers to that set of tuples or to the value(s) in the cell(s) that its tuples specify.

Syntactically, a set may be specified in a number of ways. Perhaps the most common is to list the tuples within {}. The following query uses two sets, where the rows use a set of one dimension and the columns a set of two dimensions:

```
SELECT
{ ( [Time].[2005], [Measures].[Dollar Sales] ),
  ( [Time].[Feb, 2005], [Measures].[Unit Sales] )
}
ON COLUMNS ,
{ [Product].[Tools], [Product].[Toys] } ON ROWS
FROM [Sales]
```

Whenever one or more tuples are explicitly listed, you will need to enclose them within braces. Some MDX operators and functions also return sets. The expressions that use them do not need to be enclosed in braces if the set is not being combined with more tuples, but we will usually enclose set expressions in braces for the sake of style.

Although a single member is by default a tuple of one dimension, a set that has only one tuple is not equivalent to a tuple. As far as standard MDX is concerned, the following two examples are quite different:

```
([Time].[2001 Week 1], [Product].[HyperGizmos])
{ ([Time].[2001 Week 1], [Product].[HyperGizmos]) }
```

The first of these is a tuple, and the second is a set containing that tuple. You might think it reasonable that wherever a set is called for, you can use a single tuple and it will be interpreted as a set of one. Analysis Services 2005 can make use of tuples or members in some contexts that otherwise ask for a set. For other servers, you will need to wrap the tuple in curly braces as in the second sample just given. Hence, the following is a valid query only for Analysis Services 2005 (it needs braces around the tuple to be valid otherwise):

```
SELECT
([Time].[Jun, 2005], [Geography].[Chicago, IL]) on columns
FROM [Sales]
WHERE ([Measures].[Dollar Costs])
```

The following is valid across all MDX providers:

```
SELECT
{ ([Time].[Jun, 2005], [Geography].[Chicago, IL]) } on columns
FROM [Sales]
WHERE ([Measures].[Dollar Costs])
```

Similarly, a set that happens to contain only one tuple is still considered to be a set. To use it in a context that calls for a tuple (for example, in a WHERE clause), even if you have guaranteed that it only contains one tuple, you must still employ an MDX function (such as `.Item()`) that takes a tuple from a set.

A set can also be empty, both explicitly (for example, `{ }`) and because a function returns a set that is empty for some reason.

Every tuple in a set must have the same dimensionality (that is, the dimensions represented and their order within each tuple). The following would result in an error, because the order of the dimensions changes:

```
{ ([Time].[2005], [Measures].[Dollar Sales] ),
  ([Measures].[Unit Sales], [Time].[Feb, 2005] )
}
```


However, different sets can have whatever differing dimensionality and dimension order they need in order to serve your purposes. An empty set has no dimensionality.

In addition to creating sets by explicitly coding them, a large number of functions return sets. We will explore many of them in the following chapters.

## Queries

An MDX query result is just another cube that is a transformation of the cube that is being queried. This is analogous to a standard SQL query result, which is essentially another table. The result cube can have one, two, three, four, or more axes (up to 128 in Analysis Services 2005, and up to 64 in Analysis Services 2000 and Essbase). It is also technically possible for a query to have zero axes, but it will still return a single-cell value. Each tuple in a result axis is essentially a member of the result cube.

As described earlier, each axis of the query result is composed of a set of tuples, each of which can have one or more dimensions. When multidimensional tuples end up on the axis of a query, the order in which the dimensions appear in the tuples affects the nesting order in the axis. The first dimension listed becomes the outermost dimension, the second becomes the next outermost, and so on. The last dimension is the innermost. For example, suppose that the following set was placed on the “rows” axis in a query:



```
SELECT
...
{ ([Time].[2001], [Product].[Leather Jackets]),
  ([Time].[2001], [Product].[Silk Scarves]),
  ([Time].[1997], [Product].[Leather Jackets]),
  ([Time].[1997], [Product].[Silk Scarves])
} ON ROWS
...
```

In this case, the expected presentation for data brought back to a client through OLE DB for OLAP or ADO is shown in Figure 1-10. Note that the layout shown in Figure 1-10 is simply conventional; your applications may do something different with the results.

Queries with Zero Axes

We have mentioned twice that a query may have zero axes without really describing what that means. Two examples of zero-axis queries are

```
SELECT FROM SalesCube
```

and

```
SELECT FROM SalesCube
WHERE ([Time].[2004], [Geography].[Quebec],
      [Product].[Snorkels], [Channel].[Superstores])
```

Because no members were assigned to any (non-slicer) axis in either query, the results are considered to have zero axes and by convention would be single unlabeled cells, or at least cells with no distinct row or column headers. In all APIs that currently support MDX, the slicer information is returned as part of the query result. Whether you consider the results here to be zero-dimensional depends on whether or not you choose to ignore the dimensional information conveyed in the slicer information returned with the cell.

2001	Leather Jackets
2001	Silk Scarves
1997	Leather Jackets
1997	Silk Scarves

Figure 1-10 Typical expected client data layout.

## Axis-Only Queries

Note that all MDX queries return cells. However, many useful queries ask “What members belong in this set?” where the result that is of real interest is not cell data, but the members that are associated with cell data or member property values. A query of this form, such as “Show me the customers that make up the top 10 percent of our revenue,” will at least implicitly prepare a set of cell values as well. (It may not actually retrieve their values due to internal optimizations, though.)

This is in contrast to SQL, which will return only the columns that you request. There are uses for queries that have no interest in cells; we will look at an example of this in Chapter 7.

## More Basic Vocabulary

Now that we have introduced tuples and sets and the earlier bits of vocabulary, we can introduce three more functions that are extremely common in the MDX you will write:

- `CrossJoin()`
- `Filter()`
- `Order()`

## CrossJoin()

In many cases, you will want to take the cross-product of members (or tuples) in two different sets (that is, specify all of their possible combinations). The `CrossJoin()` function is the most direct way of combining the two sets in this way.

The syntax is as follows:

```
CrossJoin (set1, set2)
```

For example, you may want to lay out on the columns of a query the first two quarters of 2005 and the two measures Dollar Sales and Unit Sales. You would generate this set with the following expression:

```
CrossJoin (
  { [Time].[Q1, 2005], [Time].[Q2, 2005] },
  { [Measures].[Dollar Sales], [Measures].[Unit Sales] }
)
```

You would use it in this way; the results are shown in Figure 1-11:

```
SELECT
CrossJoin (
  { [Time].[Q1, 2005], [Time].[Q2, 2005] },
  { [Measures].[Dollar Sales], [Measures].[Unit Sales] }
)
ON COLUMNS ,
{ [Product].[Tools], [Product].[Toys] } ON ROWS
FROM Sales
```

CrossJoin() only takes two sets as inputs. If you want to take the CrossJoin() of three or more sets, such as times, scenarios, and products, you can do it by nesting calls to CrossJoin(), like the following:

```
CrossJoin (
  [Time].Members,
  CrossJoin (
    [Scenario].Members,
    [Product].Members
  )
)

CrossJoin (
  CrossJoin(
    [Time].Members,
    [Scenario].Members
  ),
  [Product].Members
)
```

Notice that each of these results in a set whose dimensionality is, in order: time, scenario, product. While you may have a personal preference for one method or another, when the sets are large, you may want to look out for performance differences between them in your MDX provider.

	Q1, 2005	Q1, 2005	Q2, 2005	Q2, 2005
	Dollar Sales	Unit Sales	Dollar Sales	Unit Sales
Tools	183,042.00	7,179	199,809.20	7,912
Toys	149,435.30	5,878	169,513.10	6,476

**Figure 1-11** CrossJoined dimensions on columns.



`CrossJoin()` is standard MDX. Microsoft Analysis Services also provides an extension to express this as “set multiplication” by using `*` (asterisk):

```
{ [Time].Members } * { [Scenario].Members } * { [Product].Members }
```

This performs the same operation as `CrossJoin()`, and may be easier for you to read and write when you are not concerned with portability.

A common use for `CrossJoin()` is to combine a single member of one dimension with a set of members on another dimension, such as creating a set in which a particular measure is referred to over a set of tuples from other dimensions. When the formula of one calculated measure involves the count of nonempty cells for another measure, this construct is required. Although it might seem preferable, you cannot construct tuples on multiple dimensions by using range operators. For example, to express the range “toothpaste in stores 1 through 10,” you might want to write something like the following:

```
([Product].[Toothpaste],  
{[Geography].[Store 1] : [Geography].[Store 10] })
```

Instead, you will need to use `CrossJoin()` (or the `*` variant), such as in the following:

```
CrossJoin (  
  { [Product].[Toothpaste] },  
  [Geography].[Store 1] : [Geography].[Store 10]  
)
```

In the phrasing in the `CrossJoin()` example, we did not use curly braces around the set; they were not needed there. However, since the function requires a set, we did use them around the single member `[Toothpaste]`, so we could convert the tuple to a set.

**NOTE** In Analysis Services 2005, `CrossJoin()` doesn’t always create a full Cartesian product. See Chapter 4 for more details.

## Filter()

Operators like `CrossJoin()` and `:` help you construct sets. In contrast, `Filter()` lets you reduce a set by including in the resulting set only those elements that meet some criteria. `Filter()` takes one set and one Boolean expression as its arguments and returns that subset where the Boolean expression is true.

The syntax for `Filter()` is:

```
Filter (set, boolean-expression)
```

For example, the expression

```
Filter (  
  { [Product].[Product Category].Members },  
  [Measures].[Dollar Sales] >= 500  
)
```

will return a set of all product category members in which the associated Dollar Sales value was at least 500. This is the first time we have used comparisons. Any Boolean expression may be used to filter the set. As a more complex example, the expression


```
Filter (  
  { [Product].[Product Category].Members },  
  ([Measures].[Dollar Sales] >= 1.2 * [Measures].[Dollar Costs])  
  AND [Measures].[Dollar Sales] >= 150  
)
```

will return a set of all product category members in which the associated sales measure value was at least 1.2 times the associated cost measure value and the sales value was at least 150. How do you specify the associated values? The answer to that has important details, which we will address in Chapter 4. For the moment, rely on your intuition.

`Filter()` works on general sets of tuples, not just on sets of one dimension's members. The following expression returns the set of all (product category, city) tuples in which the associated sales value was at least 500:

```
Filter (  
  CrossJoin (  
    [Product].[Product Category].Members,  
    [Store].[City].Members  
  ),  
  [Measures].[Dollar Sales] >= 500  
)
```

In determining the value of sales associated with each product category, or each (product category, city) tuple, you must take into account the other dimensions that are associated with sales values. For example, the first two `Filter()` expressions and the last one did not account for the time or times with which the sales values were associated. You can specify any additional dimensions' members that you need to in either the Boolean condition or in



the set. For example, if you wanted to filter based on 2000's sales in Baton Rouge, you would simply say:

```
Filter (
    [Product].[Product Category].Members,
    ([Measures].[Dollar Sales], [Time].[2000],
     [Store].[Baton Rouge, LA]) >= 500
)
```

Within the filtering operation, the cell value will be taken from the 2000 Baton Rouge sales at each product category. The result is a set of product category members.

On the more advanced side, you can also specify more members in the set. For example, the preceding operation could be specified as follows:

```
Filter (
    CrossJoin (
        {([Time].[2000], [Store].[Baton Rouge, LA]) },
        [Product].[Product Category].Members
    ),
    [Measures].[Dollar Sales] >= 500
)
```

This expression filters a set of tuples that all include 2000 and Baton Rouge, thus fixing on the correct time and store. However, the set returned would consist of tuples with dimensionality:

```
([Time], [Store], [Product]).
```

These `Filter()` expressions have introduced the concept of query context (the relevant members for the dimensions that are not part of the filter condition or the set being filtered). Every MDX expression ultimately operates in a context that is set up outside it. Nested MDX operations are resolved within the context of the operation that invokes the nested operation. We'll defer discussion of contexts here; Chapter 4 explains query evaluation and context in far more detail.

## POWER OF TUPLES

**The expression** `([Measures].[Dollar Sales], [Time].[2000], [Store].[Baton Rouge, LA]) >= 500` **is the first use of the power of tuple references to precisely refer to cells holding data values of interest. We will show how to exploit this many times; it is one of the powerful aspects of MDX.**

## Order()

To put the tuples in a set into a sequence based on associated data values, we need to use the `Order()` function.

The syntax for the `Order()` function is:

```
Order (set1, expression [, ASC | DESC | BASC | BDESC])
```

`Order()` takes a set, a criterion for ordering the set, and, optionally, a flag that indicates what sorting principle to use (ascending or descending, including or ignoring hierarchical relationships between the tuples). `Order()` returns a set that consists of the original set's tuples placed in the new order. The precise operations of the orderings that include hierarchical relationships are fairly complex. Appendix A includes a complete description. Here, we will use simpler examples that don't show as much complexity.

For example, given the set of product categories in our database, we may want to sort them in descending order by sales. A very simple query for this may look like the following; its results are shown in Figure 1-12:

```
SELECT
{ [Measures].[Dollar Sales] } on columns,
Order (
  [Product].[Product Category].Members,
  [Measures].[Dollar Sales],
  BDESC
)
on rows
FROM [Sales]
WHERE [Time].[2004]
```

Often, you will want to use particular tuples to precisely specify the sort criteria. For example, let's say that you want to sort specifically by the profit realized in 2005 by all customers. This would be expressed by the following:

```
Order (
  [Product].[Product Category].Members,
  ([Measures].[Profit], [Time].[2005], [Customer].[All Customers]),
  BDESC
)
```

For example, the following query (whose results are shown in Figure 1-13) provides multiple time periods, customers, and measures, so you would want to use a tuple to pick out which ones to sort on:

```

SELECT
CrossJoin (
  {[Time].[2004], [Time].[2005]},
  CrossJoin (
    { [Customer].[Northeast], [Customer].[West] },
    { ([Measures].[Dollar Sales], [Measures].[Unit Sales] ) }
  )
) on columns,
Order (
  [Product].[Product Category].Members,
  ([Measures].[Unit Sales], [Time].[2005],
  [Customer].[All Customers]),
  BDESC
) on rows
FROM [Sales]

```

As you can see, the rows are sorted by the values related to the (2005, West, Dollar Sales) tuple, which corresponds to the next-to-last column.

Because `Order()` works on tuples, you can also sort the interesting (product and store) combinations by their unit sales. For example, the following expression filters (product and promotion) tuples by Dollar Sales, then orders each resulting (product and promotion) tuple according to its unit sales in 2005, and returns them; Figure 1-14 shows the resulting order.

```

Order (
  Filter(
    CrossJoin(
      [Product].[Product Category].Members
      , [Promotion].[Media].Members
    )
    , [Measures].[Dollar Sales] >= 500
  )
  , ([Measures].[Unit Sales], [Time].[2005])
  , BDESC
)

```

	Dollar Sales
Tools	894,495.80
Computers,Peripherals	847,595.00
Toys	768,126.20
Camping, Hiking	646,902.40
Phones	640,023.80
Outdoor Gear	572,151.00
Sports Equipment	541,098.50
Exercise, Fitness	534,700.60
TV, DVD, Video	500,910.80

**Figure 1-12** Result of query using `Order()`.

	2004	2004	2004	2004	2005	2005	2005	2005
	Northeast	Northeast	West	West	Northeast	Northeast	West	West
	Dollar Sales	Unit Sales	Dollar Sales	Unit Sales	Dollar Sales	Unit Sales	Dollar Sales	Unit Sales
Computers, Peripheral	118,438.80	4,612	147,504.80	5,791	148,902.90	5,695	196,868.70	7,318
Tools	132,914.20	5,282	125,364.10	5,269	152,571.30	5,916	176,282.70	6,850
Toys	103,469.60	4,220	108,991.00	4,454	106,096.30	4,233	141,530.10	5,370
Phones	120,098.40	4,913	98,190.80	3,871	125,115.60	4,787	131,831.80	4,869
Outdoor Gear	73,977.90	3,185	75,927.10	2,962	84,289.00	3,193	128,668.20	4,892
Camping, Hiking	147,407.10	5,863	69,414.90	2,819	170,091.00	6,467	123,632.60	4,727
Personal Care	31,457.50	1,312	85,282.60	3,328	53,120.30	2,083	117,788.00	4,371

**Figure 1-13** Result of more complex Order() query.

Electronics	Radio
Outdoor & Sporting	Radio
Electronics	Boat
Electronics	newsp
Electronics	email
Outdoor & Sporting	newsp
Electronics	Personal

**Figure 1-14** Result of sorting tuples.

Note that the BDESC variant breaks (that is, ignores) the hierarchy. You'd get back a more complex and possibly more interesting ordered set if you instead chose DESC, which respects the hierarchy and the dimensional components of tuples. See the section "Ordering Sets" in the description of Order() in Appendix A for a full description of ordering sets.

## Querying for Member Properties

An MDX query can also retrieve member properties defined for members in a cube along the query's axes using the DIMENSION PROPERTIES clause. For example, the following will query for the zip code and hair color of customers returned in the query:

```
SELECT
  { [Customer].[Akron, OH].Children }
  DIMENSION PROPERTIES [Customer].[Zip Code],
    [Customer].[Individual].[Hair Color]
on columns,
  { [Product].[Category].Members } on rows
FROM Sales
WHERE ([Measures].[Units Sold], [Time].[July 3, 2005])
```

The `DIMENSION PROPERTIES` clause comes between the set for the axis and the `ON AXIS (n)` clause. The `DIMENSION` keyword is optional; you could also write:

```
SELECT
  { [Customer].[Akron, OH].Children }
  PROPERTIES [Customer].[Zip Code],
               [Customer].[Individual].[Hair Color]
ON COLUMNS,
...
```

Member properties can be requested on any axis but not on the slicer in standard MDX. (Essbase 9 extends MDX to allow a `DIMENSION PROPERTIES` clause on the slicer.)

If a member property is defined for only a single level of a dimension, and the query axis includes multiple levels, the query will succeed and you will just not get values for members at the inapplicable levels. If a member is repeated in a result axis, its related member property value will be repeated too.

Properties can be identified either by using the name of the dimension and the name of the property, as with the zip code property just given, or by using the unique name of the dimension's level and the name of the property, as with the hair color property.

**NOTE** While the values of properties requested with the `PROPERTIES` statement in an MDX query are returned, along with all other result information, in the result data structures, it is up to the client application to retrieve and utilize this information.

Both intrinsic and database-defined member properties can be queried. Analysis Services defines intrinsic properties named `KEY`, `NAME`, and `ID`, and every level of every dimension has them. (These, and other kinds of properties for Analysis Services, are described in Appendix C.) For example, the `KEY` property of a Product dimension's SKU level is named `[Product].[SKU].[KEY]`. The member key property contains the values of the member keys as represented in the dimension table. The member name property contains the values of the member names as represented in the dimension table. The `ID` property contains the internal member number of that member in the dimension-wide database ordering. (Since these properties are maintained internally, your application should not use them to avoid problems with ambiguous names.) Essbase 9 defines intrinsic properties named `MEMBER_NAME`, `MEMBER_ALIAS`, `GEN_NUMBER`, `LEVEL_NUMBER`, `MEMBER_UNIQUE_NAME`, `IS_EXPENSE`, `COMMENTS`, and `RELATIONAL_DESCENDANTS`. These are the names for which Essbase is case-sensitive: if you enclose these names in `[]`, they should be in all caps (for example, `[MEMBER_NAME]`).

When property names between levels of a dimension are ambiguous, you can get ambiguous results if you query for member properties on the axis of a query. For example, every layer of an organizational dimension may have a `Manager` property for each member above the leaf. Consider the following query fragment:

```
SELECT { Descendants ([Organization].[All Organization],
[Organization].[Junior Staff], SELF_AND_ABOVE )
PROPERTIES [Organization].[Manager] on columns
... }
```

When the query is executed, it will return the specific `Manager` property for only one level. It is not a good idea to rely on whatever level that would happen to be. (In our experience, it would be the lowest level in the query, or the `Junior Staff` level in this case.) Members belonging to that level will have a valid `[Manager]` value; members belonging to other levels won't. Suppose that, instead, you queried for each level's properties independently, as with the following:

```
SELECT { Descendants ([Organization].[All Organization],
[Organization].[Junior Staff], SELF_AND_ABOVE )
PROPERTIES
[Organization].[Executive Suites].[Manager],
[Organization].[Middle Managers].[Manager],
[Organization].[Junior Staff].[Manager] on columns
... }
```

In this case, the property for each level at each level's member will arrive appropriately filled in (and be empty at members of the other levels). However, when you access properties in member calculations, there won't be any ambiguity. Suppose, for example, that some calculated member referred to (in Microsoft syntax) `[Organization].CurrentMember.Properties ("Manager")`. (Appendix A provides a detailed reference to this function, and we also use it in the "Using Member Properties in MDX Expressions" section of Chapter 3.) The lookup of this value is done on a cell-by-cell basis, and at each cell the particular manager is unambiguous (though the level of manager to which it refers may change). For this case, you can easily and simply reference member properties on multiple levels that share the same name.

## Querying Cell Properties

Querying for specific cell properties is fairly tightly bound to the programming layer that retrieves results from a query. In keeping with the nonprogramming focus of this book, we won't cover all of the programming details



here. (We do introduce client programming in Chapter 15.) However, we will explain the basic model that querying for specific cell properties supports and how an application might use it. This explanation is relevant to OLE DB for OLAP, ADO MD and ADO MD.Net, and XMLA.

Every query is a specification of one or more result cells. Much as each member is able to have one or more related properties, each result cell also has more than one possible result property. If a query specifies no cell properties, then three properties are returned by default: an ordinal number that represents the index of the cell in the result set, the raw value for the cell in whatever data type is appropriate, and the formatted textual value for the cell. If the query specifies particular cell properties, then only those are returned to the client. We discuss formatting the raw value into text in the section “Precedence of Display Formatting” in Chapter 4. The ordinal cell index value is germane to client tools that are querying the data that has been generated through OLE DB for OLAP or XMLA. Other cell properties can be queried for, which can be specified for any measure or calculated member in the cube. The full list of cell properties and how they are used in OLE DB for OLAP and ADO MD/.Net is found in Appendix C.

The way to specify cell properties in a query is to follow the slicer (if any) with the `CELL PROPERTIES` keywords and the names of the cell properties. For example, the following query

```
SELECT
{ [Measures].[Units Returned], [Measures].[Value Returned] } on columns,
{ [Time].[2000], [Time].[2001] } on rows
FROM InventoryCube
CELL PROPERTIES FORMATTED_VALUE
```

returns to the client only the formatted text strings that correspond to the query results. Generally speaking, clients that render their results as text strings (such as spreadsheet-style report grids in ASP pages) will be most interested in the formatted values. Clients that render their results graphically (such as in bar charts where each height of each bar represents the value of the measure at that intersection) will be most interested in the raw values. An Excel-based client will benefit from retrieving the raw value and the format string together. Other OLE DB for OLAP standard properties available in Analysis Services enable string formatting, font, and color information to be stored and retrieved for measures and calculated members. This gives you server-side control over useful client rendering operations. Analysis Services adds more cell properties covering cell calculation specifics.

Our discussion of `CREATE MEMBER` in Chapter 12 describes how to specify the various cell properties that should be associated with calculated members. In Chapter 4, we describe how calculated members influence cell properties in queries.

**TIP** Some development teams avoid using formatted values, claiming that it is preferable to separate formatting from the database. However, that requires mid-tier or front-end code to be conscious of the meanings of individual members and tuples, making generic display code quite complicated, while not delivering one jot of different capability. You can also think of that as breaking encapsulation. In our experience, they are extremely helpful, and can always be overridden by client code anyway if you really want to.

## Client Result Data Layout

Although the result of an MDX query is a data structure, it is typically converted to some human-accessible rendering. When more than one dimension is on an axis, there is a convention that you should be aware of for how clients use the results. Recall that the order of dimensions in a tuple is important in MDX. Consider, again, the sample query involving `CrossJoin()`:

```
SELECT
CrossJoin (
  { [Time].[Q1, 2005], [Time].[Q2, 2005] },
  { [Measures].[Dollar Sales], [Measures].[Unit Sales] }
)
ON COLUMNS ,
{ [Product].[Tools], [Product].[Toys] } ON ROWS
FROM Sales
```

The Measures dimension is the second dimension in each of the tuples, and the Time dimension the first. The first dimension in a tuple is the one at the outermost level of nesting as it comes out of the `CrossJoin()`, and it is also the outermost one in a client’s rendering of the results. The last dimension in a tuple is the innermost, both from `CrossJoin()` and in client’s rendering.

The order of dimensions in tuples, when they show up in an axis, is reflected in the order of the dimensions in the returned data structures.

	Q1, 2005	Q1, 2005	Q2, 2005	Q2, 2005
	Dollar Sales	Unit Sales	Dollar Sales	Unit Sales
Tools	183,042.00	7,179	199,809.20	7,912
Toys	149,435.30	5,878	169,513.10	6,476

**Figure 1-15** Typical client rendering of rows and columns.



## Summary

---


While we have covered only the basics of MDX queries, we really have covered quite a lot, and fairly quickly. Let's summarize what we have gone over:

- The SELECT . . . FROM . . . WHERE framework of a query
- Axis names and numbering, and the slicer
- The tuple and set data model
- The most basic and frequent functions and operators: braces ({}), commas (,), parentheses for tuple construction (()), and the colon (:) for range construction; .Members, .Children, Descendants(), CrossJoin(), Filter(), and Order()
- Referencing member properties as additional query results
- Referencing cell properties
- Including comments
- Removing empty slices with NON EMPTY

We have also tried to emphasize the modular nature of MDX and explain expressions as an independent concept from queries. Whew! You might want to take a break before plunging on to the next chapter. In the next chapter, we will build on this understanding to start performing calculations.




—



—

|



|

# Introduction to MDX Calculated Members and Named Sets

In addition to selecting data, MDX provides the ability to define new calculations, and to define named sets of tuples that can be referred to in other selections and functions. These can be done within a query, which frees the database developer from needing to account for all applications of the database, and it frees the users and/or applications from relying on the database for all calculations and set or filter definitions. In this chapter, we introduce you to creating calculations using the MDX construct called “calculated members.” Calculated members are a form of DDL (*data definition language*) that servers such as Analysis Services and Essbase allow to be defined as part of the database as well as part of a query. The way they behave is the same in either case, so we will look at how they behave from the point of view of queries. We will also introduce you to creating named sets.

Servers such as Analysis Services and Essbase also allow MDX-based calculations to be defined at the database, and Analysis Services allows them to be defined within a client session as well. Analysis Services also allows named sets to be defined at the server. Essbase 9 provides an alternate technique to accomplish the same result within MDX queries, which will be discussed in that section.

## Dimensional Calculations As Calculated Members

---


If you have some experience with SQL, you will find that the numerical calculations that are straightforward in SQL are also straightforward in MDX. However, a great many calculations that are very difficult in SQL are also straightforward in MDX. If you have experience with spreadsheet formulas (with their absolute and relative cell references in two or three dimensions), then you are already familiar with some of the basic concepts of dimensional calculations. MDX provides a much clearer language in which to create calculations, however, and you don't have to put a formula into each cell to calculate the cell. Rather, you can specify a formula and control the range of cells that all share that formula.

Not all calculations need to be put into calculated members. Only those whose results you wish to return as values for cells to a client application need to be put in calculated members. When a calculation is needed somewhere else (for example, as a basis for selecting or sorting members), an arithmetic expression will do just as well. However, you may wish to use a calculated member to package the logic of the calculation if you are going to use it more than once.

The simplest form of a calculated member is a named member in some dimension, and a formula with which to calculate cells related to that member. Calculated members exist to provide a place to put the results of calculations. Every query to a cube results in a collection of cells, where every cell is identified by one member from each dimension of the cube. There are no members without names, and there are no areas of the query that fall outside the cube. Therefore, we add a member to add a slot to a dimension that will hold the results.

**TIP** Whenever you are trying to perform a calculation in MDX, you need to define a calculated member (or perhaps some other calculated entity in Analysis Services) to contain it and then reference this entity in a query.

Under the hood of the Analysis Services 2005 BI Development Studio and the Analysis Services 2000 Analysis Manager, the user interface generates exactly the same MDX statements that create the calculated members that will be discussed in the remainder of this chapter. So, although the calculated members are metadata entities, they are also MDX language constructs. Two of the purposes of this chapter are to give you an understanding of MDX queries and of how to create calculated members for cubes in databases.



**CALCULATIONS IN ANALYSIS SERVICES**

Calculated members are only one of six different ways to compute the value of a cell in Microsoft Analysis Services 2000, and four of those ways involve MDX. Analysis Services 2005 adds yet another MDX-based way. This section focuses on calculated members because they are a workhorse, and one of only two ways that an OLAP client can define without special permissions. The concepts you gain in understanding how calculated members work go a long way towards helping you understand how to use the other techniques.

## Calculated Member Scopes

The MDX specification allows three different lifetime scopes for calculated members:

- Within the duration of a query (“query scope”)
- Within the duration of a user session (“session scope”)
- Persistent within a cube until the calculated member is dropped (“global scope”)

Analysis Services supports the first two of these scopes. Essbase supports the first one in the syntax we discuss in this chapter; calculated members defined at the server are neither created nor dropped with DDL but rather via outline editors and API calls.

## Calculated Members and WITH Sections in Queries

In order to introduce calculated members, we also need to introduce a new and optional part of a query. As this section begins with the keyword `WITH`, we will refer to this part of the query as the “`WITH` section.”

The `WITH` section of a query comes before the `SELECT` keyword and forms a section where the definitions private to the query are made. Calculated members and named sets (described later on) are the only two standard things that may be specified in the `WITH` section. (In Analysis Server, declaring a cell calculation and a cache of data to be loaded would be the other two things. We will cover cell calculations in Chapter 12. Caches are purely a physical optimization, which we will discuss in Chapter 14.)

For example, the following query whose results are shown in Figure 2-1 will augment the database measures of Unit Sales and Dollar Sales with a measure named [Avg Sales Price]:

```
WITH
MEMBER [Measures].[Avg Sales Price] AS
'[Measures].[Dollar Sales] / [Measures].[Unit Sales]'
SELECT
{ [Measures].[Dollar Sales], [Measures].[Unit Sales],
  [Measures].[Avg Sales Price]
}
on columns,
{ [Time].[Q1, 2005], [Time].[Q2, 2005] }
on rows
FROM Sales
WHERE ([Customer].[MA])
```

As you can see, the Avg Sales Price calculation is carried out across each of the time members, for the stated customer (and all other dimensions as well).


Let's step back from this example and identify some general features. The core syntax for defining a calculated member on a dimension is

```
MEMBER MemberIdentifier AS 'member-formula' [, properties...]
```

The three main parts are

- Member identifier, which specifies the name, dimension and hierarchical positioning of the calculated member.
- Formula, which specifies how the calculated member's results are derived.
- Optional properties, which can provide additional calculation, display, and other kinds of information.

The member must be associated with a dimension of a cube, so *Member Identifier* must contain a dimension name as a component. If the calculated member is to appear as the child of another member, that parent member's name must also appear as part of the member identifier. For example, the name `[Measures].[Ratios].[Avg Sales Price]` would make the name appear to be a child of the Ratios member in the measures dimension (although, as we mentioned in Chapter 1, you would need to use the `.AllMembers` or `AddCalculatedMembers()` functions to see the member appear in the list of children). Note that neither Analysis Services or Essbase allows calculated members to be the children of other calculated members.





**WHY DEFINE A CALCULATED MEMBER ONLY WITHIN A QUERY?**

Some kinds of calculations can only be performed by calculated members in a **WITH** clause. In particular, any calculation on members that are picked by a user in the course of a session cannot be stored in the database. We'll see examples of this later on.

	Dollar Sales	Unit Sales	Avg Sales Price
Q1, 2005	96,949.10	3,866	25.1
Q2, 2005	104,510.20	4,125	25.3

**Figure 2-1** Result of query with calculated member.

The formula is an expression that results in numbers or strings. The MDX standard does not include the quotes around the expression. To conform to the standard, the query should read:

```
WITH
MEMBER [Measures].[Avg Sales Price] AS
[Measures].[Dollar Sales] / [Measures].[Unit Sales]
SELECT ...
```

Since the first two versions of Microsoft's server required quotes around the expression, most vendors support it. At the time of this writing, only Analysis Services 2005 supports the standard, but we hope that other servers do shortly as well.

MDX provides two variations on this core syntax to define calculated members. If you put this definition in a **WITH** section, the calculated member is part of the query but not available in any other query. The other variation defines a calculated member that will be available to more than one query; we will discuss that one later on.

Calculated members can be on any dimension, so you can also query for the growth in dollar and unit sales between Q1 and Q2 of 2005 with the following query (its results are shown in Figure 2-2 with the growth calculation shaded in):

```
WITH
MEMBER [Time].[Q1 to Q2 Growth] AS
'[Time].[Q2, 2005] - [Time].[Q1, 2005]'
SELECT
{ [Measures].[Dollar Sales], [Measures].[Unit Sales] }
on columns,
```

```
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q1 to Q2 Growth] }
on rows
FROM Sales
WHERE ([Customer].[MA])
```

Two things stand out here: First, it is just as easy to calculate new rows as it is new columns. If you are familiar with SQL, you may appreciate the uniformity of the syntax for this. Second, we have the same syntax for measures, time, and any other dimensions that you want to define a calculation in.

Formula Precedence (Solve Order)

So far, we have only considered formulas for members of one dimension. You will very likely have formulas on members of more than one dimension, which raises the issue of what you should do when these formulas intersect. For example, consider the set of base and calculated cells shown in Figure 2-3. They are combined from the queries for Figures 2-1 and 2-2, where each calculated slice has numbers in italics.

In the example in Figure 2-3, there are formulas in two different dimensions; the cell in which they overlap shows two possible values. This cell has two possible formulas: (Q2 Avg Sales Price - Q1 Avg Sales Price) or (Difference in Sales / Difference in Units).

	Dollar Sales	Unit Sales
Q1, 2005	96,949.10	3,866
Q2, 2005	104,510.20	4,125
Q1 to Q2 Growth	7,561.10	259

Figure 2-2 Query result with calculated time growth member.

	Dollar Sales	Unit Sales	Avg Sales Price
Q1, 2005	96949.1	3866	25.1
Q2, 2005	104510.2	4125	25.3
Q1 to Q2 Growth	7561.1	259	0.3 or 29.19

Figure 2-3 3 x 3 cells with formulas: a ratio and a sum.

How do you control the ordering of calculations among dimensions? This issue is sometimes called dimensional precedence, or formula overlap.

The mechanism that standard MDX uses for dealing with dimensional formula precedence is called *solve order*. Every calculated member has an associated solve order number, which is an integer that says what the calculation priority of the member is. When calculated members overlap on a cell, the member with the highest solve order number “wins” and is used to calculate the cell. For example, the following query uses solve order numbers to make sure that the difference of average sales prices is defined as the difference of the ratios, not the ratio of the differences (resulting in the value 0.3 in the intersecting cell):

```
WITH
MEMBER [Measures].[Avg Sales Price] AS
'[Measures].[Dollar Sales] / [Measures].[Unit Sales]',
SOLVE_ORDER = 0
MEMBER [Time].[Q1 to Q2 Growth] AS
'[Time].[Q2, 2005] - [Time].[Q1, 2005]',
SOLVE_ORDER = 1
SELECT
{ [Measures].[Dollar Sales], [Measures].[Unit Sales],
  [Measures].[Avg Sales Price]
}
on columns,
{ [Time].[Q1, 2005], [Time].[Q2, 2005], [Time].[Q1 to Q2 Growth] }
on rows
FROM [Sales]
WHERE ([Customer].[MA])
```

The standard behavior is that if you don't specify a number when you specify the formula for the member, it defaults to 0. The number is fixed when you create the calculation and cannot be changed without dropping and recreating it. The numbers are simply relative precedence numbers, so there is no real requirement that the smallest number you use be 0. Nor is there any requirement that you use 2 or 1 if the highest number in use is 3 and the lowest is 0. See Chapter 4 for a detailed discussion on solve orders.

Note that you don't need to care about solve order when it doesn't make a difference which dimension wins. If one dimension's formula specifies a sum or a difference and another dimension's also specifies a sum or a difference, a solve order tie between them won't change the results. Similarly, multiplication and pure division will be indifferent to solve orders.

**SOLVE ORDERS****ANALYSIS SERVICES**

**Analysis Services supports solve order numbers down to -8191 and up to 2,415,919,103. Some of these solve order numbers are reserved for particular kinds of calculations. Although Microsoft states that the concept of solve order is deprecated for Analysis Services 2005, this can only be true for server-defined calculated members; it cannot really be deprecated for calculated members defined in a session or a `WITH` clause.**

**ESSBASE**



**Essbase 7.1.3 and later supports default solve orders on a dimension-wide basis (defined at the server). Specifying a solve order of 0 actually means to use the dimension-wide solve order, which is 0 by default but could be different.**

A few paragraphs ago we referred to other parts of a calculated member that we can specify in the member definition. The solve order property shown in the preceding query is one of them.

Note the following syntactic point about defining multiple calculated members in the `WITH` section of the query. Each member definition is simply followed by the next one. A comma is used to separate the formula definition from the solve order definition, but no punctuation, such as a comma, semicolon, or other device, is used to separate the end of one from the beginning of the next. Instead, they are separated by the `MEMBER` keyword.

You should keep two other points in mind regarding solve orders. First, if members on two different dimensions have the same priority, it's up to the server to decide which one goes first, and they generally use some list ordering of dimensions in the cube as the basis. If the DBA rearranges the ordering, the MDX may behave differently.

You should only let formulas on different dimensions have the same solve order number when the formulas are commutative (either when they all involve only addition and subtraction or when they all involve only multiplication and division). Second, the solve order only affects the priority of calculation between dimensions. The database still uses actual formula dependencies to determine what to calculate first. For example, consider the following four formulas. Figure 2-4 shows these four formulas and their inputs laid out on a grid, together with the formula that is actually in use for any given cell.



**SOLVE ORDER TIES****ANALYSIS SERVICES**

A calculated member solve order tie in Analysis Services is won by the dimension that appears later in the server's list of dimensions. Measures is always the last dimension in the list, so it always effectively wins. Note there are other kinds of definitions that also have solve order and may win the tie instead—see Chapter 4 for details.

**ESSBASE**

Winning a solve order tie in Essbase is different between BSO and ASO applications. In a block-storage cube, and in ASO cubes prior to version 7.1.2, it is won by the dimension that appears earlier in the server's list of dimensions (which is generally compatible with precedence in Essbase's other languages). In 7.1.2 and after, a tie is won by the dimension that appears *later* in the list of dimensions.

```
[Measures].[Profit]
  AS '[Measures].[Sale Amount] - [Measures].[Total Cost]',
  SOLVE_ORDER = 0
[Scenario].[Amount of Variance]
  AS '[Scenario].[Actual] - [Scenario].[Planned]',
  SOLVE_ORDER = 1
[Measures].[Percentage Margin]
  AS '[Measures].[Profit] / [Measures].[Sale Amount]',
  SOLVE_ORDER = 2
[Scenario].[Percentage Variance]
  AS '[Scenario].[Amount of Variance] / [Scenario].[Planned]',
  SOLVE_ORDER = 3
```

Calculated members for a cube's dimension may be defined at the server or at the client. Calculated members defined at the server will be visible to all client sessions that can query the cube and can be used in any number of queries. A client can also define such calculated members as well. Clients and servers do this by using the second variation of the syntax for creating calculated members: the CREATE MEMBER command.

**NOTE** Essbase 9 does not support creating calculated members in a client session with CREATE MEMBER.

	Sale Amount	Total Cost	Profit	Percentage Margin
Actual				
Planned				
Amount of Variance				
Percentage of Variance				

**Figure 2-4** Map of calculated member definitions and overlap on a grid.

Calculated members defined with the `CREATE MEMBER` command must be named with the cube as well as the dimension that they are to be a part of. The `CREATE MEMBER` command is not part of a query that uses `SELECT`, but is its own statement. Other than that, the core syntax for naming the member and defining its formula and other properties is basically the same as a `WITH`-defined member. For example, the following MDX statement will create `[Scenario].[Amount of Variance]` on the `Scenario` dimension used by the `[Sales Cube]` cube:

```
CREATE MEMBER [Sales Cube].[Scenario].[Amount of Variance] AS
 '[Scenario].[Actual] - [Scenario].[Planned]', SOLVE_ORDER = 1
```

This calculated member will only be visible to queries on the `[Sales Cube]` cube. Queries to other cubes, even if they also use the scenario dimension, will not be able to see this calculated member, just like they cannot see any other information about another cube. The `CREATE MEMBER` statement defines a calculated member for a dimension that can be used by any query (until the member is dropped or the client exits), and additionally in Analysis Services will exist in the dimension's metadata visible through OLE DB for OLAP. (In Analysis Services, this metadata will only be visible on that client; that metadata will not be visible at the server or at any other client attached to that server.)

Note that the name of the cube was part of the specifications of the member to create. This is necessary because of OLAP providers such as Analysis Services 2005, a session isn't ordinarily associated with any particular cube. However, in Analysis Services 2005, you can include a cube name in the connection string or connection parameters, which you use to establish a session (see Appendix B for details). In that case, you can omit the cube name portion of a named set definition.

When a query uses calculated members in Essbase and in Analysis Services 2000, all solve order numbers from all the calculated members in the query are thrown together regardless of their source. A formula defined in the `WITH` section of a query as having `solve order = 2` will lose to a server-defined formula or a formula defined in a `CREATE MEMBER` statement as having `solve order = 3`. This allows for very graceful interaction of definitions. For example, key ratios can be defined at the server, such as Average Sales Price and Period-to-Period % Growth, clients can create calculations such as totals across custom sets of members (one of the many uses for named sets), and a query can take the correct ratio of sums.

**NOTE** Analysis Services 2005 behaves somewhat differently than this. How it handles solve orders is treated in depth in Chapter 4.

At the time a query is constructed, you can know the solve orders for all calculated members included in the cube definition on the server. However, when a database is constructed, you obviously cannot know the formulas used in queries and their solve orders. Furthermore, since solve orders are integer numbers, you cannot slip a calculated member into a query between two members whose solve orders are 1 and 2 by giving the new member a solve order of 1.5.

For these reasons, you may want to leave gaps in the solve order numbers that are used for calculated members created as part of a cube's definition (through the GUI or through ASO if you are creating them through your own code). For example, the lowest-precedence number at the server might be 10, the next one 20, and so on. If you ever programmed in classic BASIC, this procedure should be familiar (remember line numbers running 10, 20, 30, and so on?). Analysis Services' solve order numbers can run up to 2,147,483,647, so you have plenty of headroom here. While the Essbase maximum solve order number of 127 may seem limiting, you still have lots of room if you leave gaps of 5.

**NOTE** As of Essbase 7.1.2, when you define a calculated member, its solve order must be at least as large as the largest solve order of one of its inputs, or you will get an error message (instead of perhaps wrong results) calculating affected cells when the query is run. When you are building queries, you may need to be aware of the solve orders in use at the server.

The solve order of calculated members is one facet of the concept of formula application ranges. Basically, every formula that you will create will apply to some set of locations in your database. As far as MDX semantics are concerned, the formula that you define for a calculated member will be calculated over every cell in the database that intersects that member. This may or may not be what you want, depending on your circumstances. You may, at times, want some formulas to calculate differently, depending on what level they are at in a dimension. Profitability may be calculated differently in different countries, for example. A formula to compute GNP will not apply to a subnational level of data. We will explore techniques for controlling application ranges in Chapter 7. Also, Analysis Services 2005's MDX Scripts provide a different way to approach the problem, and we discuss them in depth in Chapter 13.

## Basic Calculation Functions

MDX includes a collection of functions with which to calculate numbers, which is rather small, so both Microsoft and Hyperion have added a number of useful extensions. Now that we have looked at the outline of how to define a calculation, let's examine some of the functions available.

### Arithmetic Operators

As you might expect, MDX provides support for the most common arithmetic operators:

FUNCTION OR OPERATOR	CALCULATES:
+ - * /	Addition, subtraction, multiplication, and division
-	For unary negation
()	Parentheses for grouping operations
^	For "raised-to-the-power-of" (Microsoft-only)

Addition, subtraction, multiplication, and division are expressed as usual, with the usual order of operations. You can apply grouping parentheses as well to specify the order of operations. Note that parentheses are used for constructing tuples as well, but there should be no cases where the use of the parentheses is ambiguous.



## Summary Statistical Operators

MDX also provides a collection of summary statistical operators, including:

FUNCTION OR OPERATOR	CALCULATES:	EXTENSION
Avg ( )	Average (mean) of values	
Aggregate ( )	Aggregate of values based on server definition of aggregation function	
Count ( )	Count of values or tuples	
DistinctCount ( )	Distinct count of values or tuples	Microsoft
Sum ( )	Sum of values	
Max ( )	Maximum of values	
Median ( )	Median value from a set	
Min ( )	Minimum of values	
NonEmptyCount ( )	Count of values or tuples	Hyperion
Stdev ( )	Sample standard deviation across values	
StdevP ( )	Population standard deviation across values	Microsoft
Var ( )	Sample variance across values	
VarP ( )	Population variance across values	Microsoft

With the exception of Count ( ) and NonEmptyCount ( ) , these functions all aggregate across a set of numbers and return a single number. Count ( ) and NonEmptyCount ( ) look at the presence of data values in cells rather than the values themselves.

The syntax and usage for each of these aggregation functions is described in the following sections.

## Avg()

The `Avg()` function takes a set of tuples and an optional numeric expression, and returns the average (mean) value of the expression over the set. By default, empty cells aren't counted, so its result is the sum across the set divided by the count of nonempty values for the expression across the set. Essbase extends the standard function with an optional `INCLUDEEMPTY` flag; when it is provided, the function returns the sum across the set divided by the number of tuples in the set. If the numeric expression isn't provided, it will take the average value over the cells related to the set in the current context.

Standard syntax:

```
Avg (set [, numeric_value_expression])
```

Hyperion syntax:

```
Avg (set [, numeric_value_expression] [,INCLUDEEMPTY] )
```

Example:

```
Avg (
    {[Product].[Tools],
     [Product].[Toys],
     [Product].[Indoor, Outdoor Games] },
    [Measures].[Dollar Sales]
)
```

## Count(), .Count

The `Count()` function takes a set of tuples and an optional flag to signal including or excluding tuples with empty related data values. In Analysis Services, the default behavior is to count all tuples in the set (as though `INCLUDEEMPTY` had been specified). If `EXCLUDEEMPTY` is specified, the related cells are evaluated and only the count of tuples with nonempty cells is returned. In Essbase, if the `INCLUDEEMPTY` flag is included, the function returns the count of tuples in the set, otherwise it defaults to excluding tuples related to empty data values.

Note that this function behaves differently from the other aggregating functions, in that it does not take an expression for which to count the values. This makes use of it a bit tricky. Its typical use involves `CrossJoin()`ing the measure whose values might be empty in with the set. Otherwise (as will be explained in Chapter 4), you may get an infinite recursion error.

Standard syntax:

```
Count (set [,INCLUDEEMPTY] )
```

Microsoft syntax:

```
Count (set [,EXCLUDEEMPTY | INCLUDEEMPTY] )
```

Example:

```
Count (
    CrossJoin (
        { [Measures].[Unit Sales] },
        [Product].[Tools].Children
    )
)
```

To simplify programming, Microsoft also accepts the `.Count` syntax that is familiar to VB programmers as a variation on `Count ( . . . INCLUDEEMPTY )`, which makes it more clear that it is counting tuples rather than cell values:

```
{[Product].[Tools].Children}.Count
```

or

```
[Product].[Tools].Children.Count
```

## DistinctCount() (Microsoft extension)

The `DistinctCount()` function takes a set of tuples and returns the count of distinct, nonempty tuples in the set. Note that only a calculated measure may use this function. It is equivalent to `Count ( Distinct (set), EXCLUDEEMPTY )` when the set includes a measure (as shown in the example, which is similar to the correct usage for `Count()`). When the set does not include a measure, the difference is that `Count()` will return an infinite recursion error, while `DistinctCount()` will return the full number of tuples in *set*. (See Chapter 4 for an explanation of the infinite recursion.)

Syntax:

```
DistinctCount (set)
```

Example:

```
DistinctCount (
    [Measures].[Dollar Sales] *
    {[Product].[Tools], [Product].[Toys],
     [Product].[Indoor, Outdoor Games]
    }
)
```

## Sum()

The `Sum()` function takes a set of tuples and an optional numeric expression, and returns the sum of the expression over the set. If the numeric expression isn't provided, it will sum over the cells related to the set in the current context. (We haven't talked about the current context yet—we will talk about it in detail later.)

Syntax:

```
Sum (set [, numeric_value_expression])
```

Example:

```
Sum (  
    {[Product].[Tools],  
     [Product].[Toys],  
     [Product].[Indoor, Outdoor Games] },  
    [Measures].[Dollar Sales]  
)
```

## Max()

The `Max()` function takes a set of tuples and an optional numeric expression, and returns the maximum value of the expression over the set. If the numeric expression isn't provided, it will take the maximum value over the cells related to the set in the current context.

Syntax:


```
Max (set [, numeric_value_expression])
```

Example:

```
Max (  
    {[Product].[Tools],  
     [Product].[Toys],  
     [Product].[Indoor, Outdoor Games] },  
    [Measures].[Dollar Sales]  
)
```

## Median()

The `Median()` function takes a set of tuples and an optional numeric expression, and returns the median value found (the one for which an equal number of values in the set were larger than as smaller than). If the numeric expression isn't provided, it will take the median value over the cells related to the set in the current context.



Syntax:

```
Median (set [, numeric_value_expression])
```

Example:

```
Median (
    [Product].[Tools].Children,
    [Measures].[Dollar Sales]
)
```

**NOTE** Essbase does not implement this function.

## Min()

The `Min()` function takes a set of tuples and an optional numeric expression, and returns the minimum value of the expression over the set. If the numeric expression isn't provided, it will take the minimum value over the cells related to the set in the current context.

Syntax:

```
Min (set [, numeric_value_expression])
```

Example:

```
Min (
    [Product].[Tools].Children,
    [Measures].[Dollar Sales]
)
```

## NonEmptyCount() (Hyperion extension)

To make the task of counting nonempty cells more like the other aggregation functions, Hyperion provides the `NonEmptyCount()` function, which takes a set of tuples and an optional numeric expression and returns the number of tuples in the set for which the expression is not empty. If the numeric expression isn't provided, it will take the count of nonempty tuples over the cells related to the set in the current context.

Syntax:

```
NonEmptyCount (set [, numeric_value_expression])
```

Example:

```
NonEmptyCount (
    [Product].[Tools].Children,
    [Measures].[Unit Sales]
)
```

## Stdev(), Stddev()

The `Stdev()` function takes a set of tuples and an optional numeric expression, and returns the sample standard deviation across the values found. If the numeric expression isn't provided, it will take the sample standard deviation over the cells related to the set in the current context. Analysis Services accepts `Stddev()` as a synonym for `Stdev()`.

Syntax:

```
Stdev (set [, numeric_value_expression])
```

Example:

```
Stdev (  
    [Product].[Tools].Children,  
    [Measures].[Dollar Sales]  
)
```

**NOTE** Essbase does not implement this function.

## StdevP(), StddevP() (Microsoft Extension)

The `StdevP()` function takes a set of tuples and an optional numeric expression, and returns the population standard deviation across the values found. If the numeric expression isn't provided, it will take the population standard deviation over the cells related to the set in the current context. Analysis Services accepts `StddevP()` as a synonym for `StdevP()`.

Syntax:


```
StdevP (set [, numeric_value_expression])
```

Example:

```
StdevP (  
    [Product].[Tools].Children,  
    [Measures].[Dollar Sales]  
)
```

## Var(), Variance()

The `Var()` function takes a set of tuples and an optional numeric expression, and returns the sample variance across the values found. If the numeric expression isn't provided, it will take the sample variance over the cells related to the set in the current context. Analysis Services accepts `Variance()` as a synonym for `Var()`.



Syntax:

```
Var (set [, numeric_value_expression])
```

Example:

```
Var (  
    [Product].[Tools].Children,  
    [Measures].[Dollar Sales]  
)
```

**NOTE** Essbase does not implement this function.

### VarP(), VarianceP() (Microsoft Extension)

The VarP() function takes a set of tuples and an optional numeric expression, and returns the population variance across the values found. If the numeric expression isn't provided, it will take the population variance over the cells related to the set in the current context. Analysis Services accepts VarianceP() as a synonym for VarP().

Syntax:

```
VarP (set [, numeric_value_expression])
```

Example:

```
VarP (  
    [Product].[Tools].Children,  
    [Measures].[Dollar Sales]  
)
```

### Additional Functions

There are more standard and extension statistical aggregation functions that we will list only briefly here:

FUNCTION	CALCULATES:
Correlation ()	Correlation coefficient of two data series over a set
Covariance ()	Covariance of two data series over a set
CovarianceN ()	(Microsoft) Sample covariance of two data series over a set
LinRegIntercept ()	Linear regression intercept

(continued)

(continued)

FUNCTION	CALCULATES:
LinRegPoint ()	Point on linear regression line
LinRegR2 ()	Linear regression coefficient
LinRegSlope ()	Slope of linear regression line
LinRegVariance ()	Variance of set from linear regression line

You are encouraged to peruse Appendix A for the details on these functions (and many more). Note that Essbase 9 does not support this set of statistical functions. Besides these functions, there are a great many more that you will want to use. Before you move on to the next chapter, you may find it most useful to at least skim Appendix A to see what is there.

Analysis Services provides access to many of the string and numeric functions of the VBA and Excel object libraries, and you can use the functions from these libraries as though they were built into your server. In Analysis Services 2005, it does so through the stored procedure mechanism, whereas in Analysis Services 2000, it does so through the external function (or “UDF”-*user-defined function*) mechanism. These are described more fully in Chapter 10. For example, logarithms, exponentiation, substring searches, combinatorial functions, percentiles, time conversions, duration calculations, and the like are all provided through the libraries.

Essbase provides an additional set of numeric calculation functions built-in. They include the following:

FUNCTION OR OPERATOR	CALCULATES:
Abs (num)	Absolute value of <i>num</i>
Exp (N)	<i>E</i> raised to the power <i>N</i>
Factorial (N)	Factorial of <i>N</i>
Int (num)	Number rounded down to the next lowest integer
Ln (num)	Natural logarithm
Log (num)	Logarithm, base <i>N</i>
Log10 (num)	Logarithm, base 10
Mod (num, M)	Modulus <i>M</i> (remainder from <i>num</i> divided by <i>M</i> )
Power (num, M)	Number <i>num</i> raised to power <i>M</i>
Remainder (num)	Number left over after removing integer portion from <i>num</i>
Round (num, M)	Number <i>num</i> rounded to <i>M</i> places
Truncate (num)	Number <i>num</i> with fractional part removed



## Introduction to Named Sets

Building and using sets is an important part of MDX. Sometimes, it is most convenient to attach a name to a set that you've defined so you can reference it later. Standard MDX provides *named sets* to do just that. Both Analysis Services and Essbase extend the standard with set aliases to provide names for sets in additional contexts. (Essbase further provides tuple aliases as well.) These are easy to use and frequently used. The second half of this chapter has some rather involved query requirements that get solved with only a small number of features and functions working together. It both explains how some sophisticated user requests are handled and also continues to work through constructing MDX as assemblies of small simple expressions, so it may bear reading slowly and more than once.

Named sets are just names for a particular set. They can hold any kind of members or tuples so long as the set is valid; anywhere that you can use a set in MDX, you can reference a named set.

For example, the following shows the definition of a named set and its use in two places: in a calculated member and in forming the set of an axis. The result is shown in Figure 2-5.

```
WITH SET [User Selection] AS
  '{ [Product].[Action Figures], [Product].[Dolls] }'
MEMBER [Product].[UserTotal] AS
  'Sum ( [User Selection] )'
SELECT
  { [Time].[Jan, 2005],
    [Time].[Feb, 2005]
  }
ON COLUMNS,
  { [Product].[Toys],
    [User Selection],
    [Product].[UserTotal]
  }
ON ROWS
FROM Sales
WHERE ([Measures].[Unit Sales])
```

	Jan, 2005	Feb, 2005
Toys	1,989	1,695
Action Figures	405	352
Dolls	28	30
UserTotal	433	382

**Figure 2-5** A query result incorporating a named set.

The basic syntax for defining a named set that is supported by most vendors is as follows:

```
SET Set-Identifier AS 'set-formula'
```

The syntax defined by the standard is similar, just without the single quotes around the set formula. As of this writing, Analysis Services 2005 is the only product that supports this syntax, but we hope more vendors do in the near future:

```
SET Set-Identifier AS set-formula
```

Essbase 9 supports a variation on the first version of the syntax, which we will discuss in Chapter 5, after discussing context in query execution in Chapter 4.

A named set name is just a name, and is not prefixed with any dimension. A set may contain tuples for more than one dimension, and if it's empty it won't appear to have any dimensions. There are no additional properties that you can specify for it.

Once it is defined and evaluated, its collection of tuples is fixed throughout its life.

## Named Set Scopes


Named sets are scoped just like calculated members. If you create a named set in the `WITH` part of a query, it is not visible outside of the query. If you create a named set using `CREATE SET` (as with the following), then it will be available to more than one query and you can drop it with the `DROP SET` command. The following shows examples of the standard MDX for `CREATE SET` and `DROP SET`, available in Analysis Services.

```
CREATE SET [SalesCube].[User Selection] AS  
'{ [Product].[Action Figures], [Product].[Dolls] }'
```

```
DROP SET [SalesCube].[User Selection]
```

Note that the name of the set is prefixed with the name of the cube. In Analysis Services, a session isn't ordinarily associated with any particular cube. However, in Analysis Services 2005, you can include a cube name in the connection string or connection parameters that you use to establish a session (see Appendix B for details). In that case, you can omit the cube name portion of a named set definition.

In a context like an MDX script, a local cube file, or in a DSO command in Analysis Services 2000, you can also use the token `CURRENTCUBE` instead of the cube name, since the cube context is clear.



Essbase does not follow the standard, but rather uses an optional FROM clause to name the cube in which to create the set. For example, the following statements would be equivalent to the foregoing examples of CREATE and DROP:

```
CREATE SET [User Selection] AS
'{ [Product].[Action Figures], [Product].[Dolls] }'
FROM [Sales]

DROP SET [User Selection]
FROM [Sales]
```

Note that the FROM clause is optional in the same way that it is in Essbase queries. For APIs that connect to an application and cube, if the FROM clause is omitted, then the cube is inferred from that. (In the MAXL and AAS interfaces, you must use the FROM clause, because they do not have a connection to a cube.)

Although the official MDX specification says that such sets are not visible to metadata, Analysis Services will surface them to client applications (which is much more helpful than hiding them).

**NOTE** Essbase also supports an optional WHERE clause; we will see what it is used for in Chapter 5, after we have covered query evaluation context in Chapter 4.

#### ESSBASE SUBSTITUTION VARIABLES AND NAMED SETS

Although Essbase does not support named sets as such at the server, it does support “substitution variables,” which are named snippets of arbitrary text that can be substituted in and are treated as part of the MDX expression at the point it is parsed. (They have broader application than just as a proxy for named sets, but are especially relevant in this context.) For example, in Analysis Services, you might define the following at the server:

```
CREATE SET [SalesCube].[East Region Promotion Products] AS
'{ [Product].[Action Figures], [Product].[Dolls] }'
```

**In Essbase, you might define a substitution variable named “East\_Region\_Products” with the contents:**

```
{ [Product].[Action Figures], [Product].[Dolls] }
```

**You could then execute any of the following:**

```
CREATE SET [East Region Promotion Products] AS
'&East_Region_Products'
```

```
WITH SET [East Region Promotion Products] AS
'&East_Region_Products'
SELECT &East_Region_Products on axis(1) ...
```

Named sets can be defined anywhere within a `WITH` section. If a named set references a calculated member or another named set, it needs to appear after the things it references. Conversely, any calculated member that references a named set needs to appear after the definition of the set.

As with calculated members, multiple named sets can be defined in a `WITH` section. Each set definition appears in sequence, as follows:

```
WITH
SET [User Selection 1] AS '...'
SET [User Selection 2] AS '...'
```

In general, named sets and calculated members can be interspersed as necessary. There are no commas between the definitions; the `MEMBER` and `SET` keywords are enough to signal the start of a definition.

```
WITH
SET [User Selection 1] AS '...'
SET [User Selection 2] AS '...'
MEMBER [Measures].[M1] AS '...'
SET [User Selection 3] AS '...'
MEMBER [Product].[P1] AS '...'
```

As you can see, named sets are simple to create and use. In Chapter 5, we will cover them more, and look at their cousins, named set aliases.

## Summary

---

You have looked at the basics of creating calculated members and named sets. We introduced the `WITH` section, how calculated members are made members of some dimension, and described calculated member formulas and solve orders. We also introduced some of the many functions that can be used to calculate the results that a calculated member can provide. However, all of this was really describing simple mechanics. Much of the real power of MDX comes from specifying how to get the input values to these functions (the various functions that return sets, tuples, and members). In the next chapter, you start to look at how to express many kinds of common calculations used in many kinds of applications, and spend more attention to the thought process of arriving at the right expression to calculate what you want.

